

## CS352: Cloud Computing Jan-May 2019

### Container Orchestration

Building a Fault tolerant Container Orchestrator with a load balancer, an auto scaling feature from scratch

UE16CS352 | Date- 27/04/19

SNo	Name	USN	Class/Section
1	Abhishek Narayanan	01FB16ECS016	6th sem/A
2	Abhishek Prasad	01FB16ECS017	6th sem/A
3	Abijna Rao	01FB16ECS019	6th sem/A
4	Adapa Shivani	01FB16ECS022	6th sem/A

## Introduction

This project is aimed at building a container orchestrator that can perform load balancing, is highly fault tolerant, and can perform auto-scaling out or scaling in of container within an EC2 instance of Amazon AWS.

The Load balancer forwards all incoming HTTP requests equally between all running acts container. Fault tolerance is ensured by the Health Check API which monitors the health of each container by polling each container every one second, and recreates a new healthy container at the same port if a container running on a certain port was found to be unhealthy. Auto scaling api runs at every 2 minute interval and based on the no of http requests received it scales out or scales in respectively.

## RELATED WORK

There are different existing literatures portraying different approaches to load balancing among multiple virtual machines or containers. The scheduling algorithm, which is the deciding factor in determining which request to the load balancer is to be routed to which node can vary from implementation to implementation. For example, requests could be routed by the load balancer based on various approaches such as Round Robin, Hashing, etc. to name a few. In our approach, for simplicity of implementation sake, we employ the round robin technique for load balancing or equally distributing incoming requests to all existent containers.

For convenience in implementation, our auto scaler feature just monitors the total number of requests every two minutes and scales out or scales in based on this. Existing auto scaling features usually are flexible and allow the user of the cloud platform to provide input rules in the form of monitor, triggers and actions wherein a user may instruct the cloud service provider to scale out or in based on selected parameters such as CPU Utilization, storage left on existing device, number of incoming requests, etc.

As an additional effort over and above the provided specifications, we try to implement in this project, a more generic auto scaler with user provided input parameters for monitors and triggers and actions.

## EXPERIMENTS/DESIGN

### Algorithms used for Task 1:

Algorithm Health\_check:

```
// Checks if the server is functioning normally

    if ( crash == True)

        status_code ← 500

        return status_code

    try block

        connect ← connection to database

        Insert data to database

        Select data from database

        Delete data from database

        status_code ← 200

        return status_code

    if any exception occurred in try block

        status_code ← 500

        return status_code
```

Algorithm Crash\_Server:

```
// Permanently disables an Acts container

    global crash ← True

    status_code ← 200

    return status_code
```

## Algorithms used for Task 2:

### Algorithm Run\_Load\_Balancer

Listen on Port 80

Accept requests from all hosts

### Algorithm Load\_Balancer ( request)

global port\_index, active\_ports

port\_index  $\leftarrow$  (port\_index + 1)%len(active\_ports)

new\_host\_url  $\leftarrow$  "localhost:" + active\_ports[port\_index]

url  $\leftarrow$  request.host\_url replaced by new\_host\_url

method  $\leftarrow$  request.method

headers  $\leftarrow$  {key:value in request.headers such that key!="Host"}

data  $\leftarrow$  request.get\_data()

response  $\leftarrow$  Send http request having method, headers, data to url

return response.text, response.status\_code, response.headers.items()

## Algorithms used for Task 3:

### Algorithm Fault\_Tolerance

global active\_ports

while (True)

for port in active\_ports

response  $\leftarrow$  send GET request to "[http://localhost:port/api/v1/\\_health](http://localhost:port/api/v1/_health)"

if (response.status\_code == 500)

container\_id  $\leftarrow$  fetch container id of container listening through port

stop container\_id

new\_container\_id  $\leftarrow$  create new container with port\_no  $\leftarrow$  port

sleep for 1s

## Algorithms used for Task 4:

### Algorithm Auto\_scaling

```
global active_ports, total_no_of_counts

Wait till first request is received

Start timer

end_request_count ← 0

while (True):

    start_request_count ← total_no_of_counts

    no_of_required_containers=(start_request_count - end_request_count )//20 + 1

    required_container_ports ← [8000+i for i in
range(no_of_required_containers)]

    if no_of_required_containers > len(active_ports)

        for each container_port in required_container_ports

            if container_port not in active_ports

                create new container with port ← container_port

    if no_of_containers < len(active_ports)

        for each container_port in active_ports

            if container_port not in required_container_ports

                remove container with port ← container_port

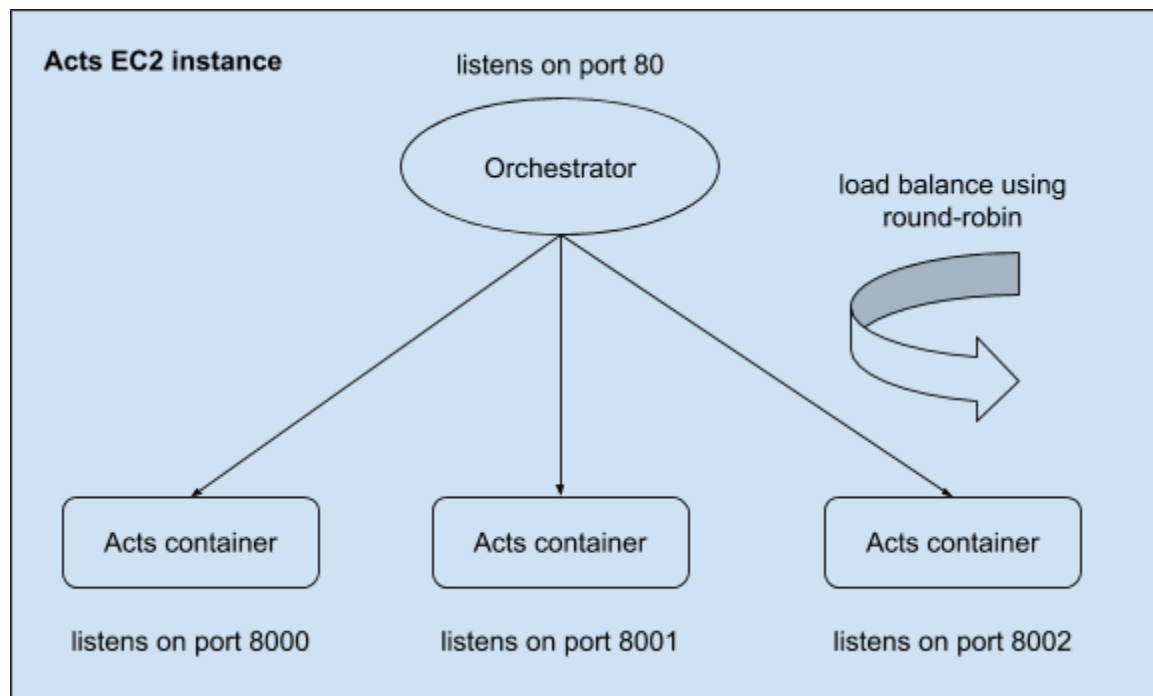
    end_request_count ← start_request_count

    sleep for 2 mins
```

### ADDITIONAL EXTENSIONS ABOVE GIVEN SPECIFICATIONS :

We extend the auto scaling feature to be more generic, such that the scaling in and scaling out of containers occurs based on user defined rules specifying how long the number of requests parameter should be monitored, what the threshold limit for every t seconds interval should be to trigger an action (which is scale in or out) and also the scale factor by which we should scale out our containers or scale in.

Further extensions in this regard would be focussed at allowing more flexibility with the parameter to be monitored as well based on user input rather than us considering only number of requests for triggering a scaling action.



The above figure pictorially portrays the architecture of our Acts EC2 setup. In the users EC2 instance, we have only one container running with all requests related to `/api/v1/users` url being routed to it by the Amazon AWS Load Balancer.

## TESTING/RESULTS

The features involved in the project, namely the load balancer, the health monitoring api for ensuring fault tolerance and auto scaler have been tested manually by writing a test script sending post requests for users and categories and 170 post requests for acts.

The test script for evaluation was also run locally and our cloud backend was robust enough to pass all required test cases, the screenshot of which is attached below :

## 16CS352: Cloud Computing – Final Project

Evaluation for CC\_016\_017\_019\_022

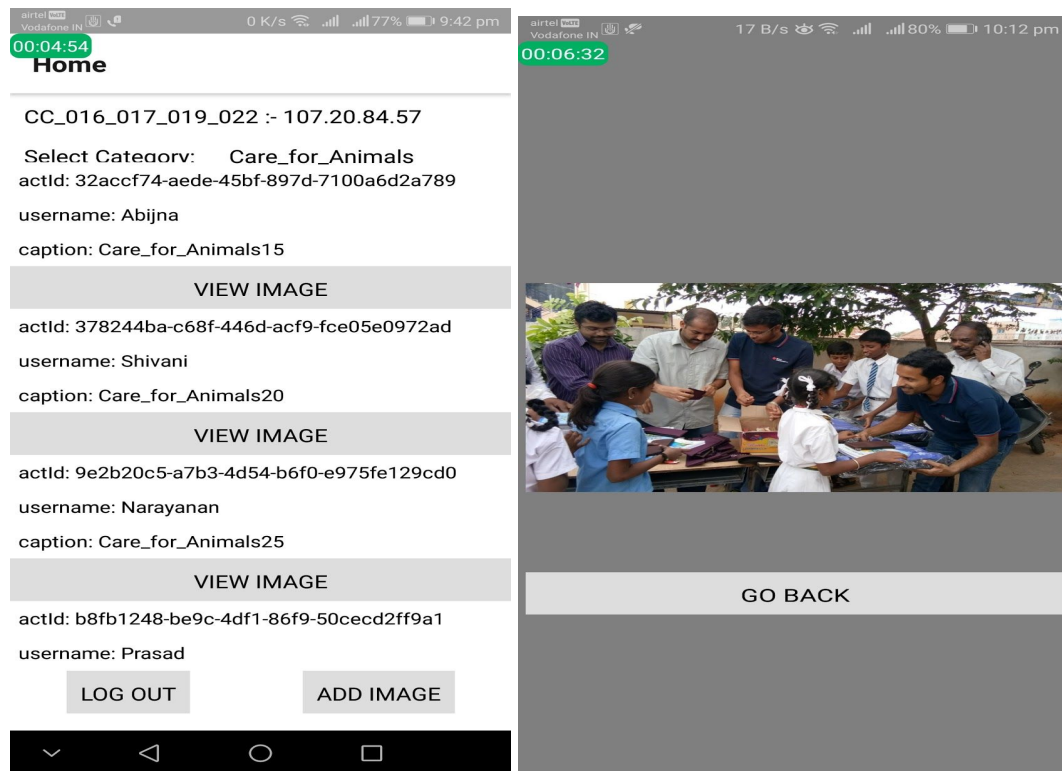
SelfieLess IP address: <http://Load-Balancer-1712257144.us-east-1.elb.amazonaws.com>

Acts-only IP address: <http://107.20.84.57>

Acts-only EC2 Username: ubuntu

1. One Acts container running with ID 25816b674570
2. Acts container 25816b674570 is listening on port 8000
3. Acts container 25816b674570 is healthy
4. 20 successful API requests made to orchestrator
5. Script slept for 2 minutes
6. Second Acts container now running with ID ac0826128513
7. Acts container ac0826128513 is listening on port 8001
8. 6 successful API requests made to orchestrator
9. API requests evenly distributed b/w 2 Acts containers in round-robin manner
10. Successfully crashed container 25816b674570 with `/api/v1/_crash`
11. Replacement Acts container running with ID fbcdb542aecb
12. Replacement Acts container fbcdb542aecb is listening on port 8000
13. Replacement Acts container fbcdb542aecb is healthy
14. Script slept for 2 minutes
15. Orchestrator successfully scaled down containers. Only 1 Acts container running with ID fbcdb542aecb

Apart from this, we have also integrated the SelfieLessActs Mobile App with our cloud backend, the screenshots of which are portrayed below :



## REFERENCES

The references referred to for achieving the tasks in the project are as follows ::

<https://docs.docker.com/develop/sdk/examples/>

<https://runnable.com/docker/python/dockerize-your-python-application>

<http://containertutorials.com/py/docker-py.html>

<https://github.com/docker/docker-py>

<https://docker-py.readthedocs.io/en/stable/containers.html>

[https://www.bogotobogo.com/python/Multithread/python\\_multithreading\\_Synchronization\\_Lock\\_Objects\\_Acquire\\_Release.php](https://www.bogotobogo.com/python/Multithread/python_multithreading_Synchronization_Lock_Objects_Acquire_Release.php)

<https://www.geeksforgeeks.org/multithreading-in-python-set-2-synchronization/>

<https://hackernoon.com/synchronization-primitives-in-python-564f89fee732>

<https://stackoverflow.com/questions/37144357/link-containers-with-the-docker-python-api>

<https://www.laurentluce.com/posts/python-threads-synchronization-locks-rlocks-semaphores-conditions-events-and-queues/>

## EVALUATIONS (Leave this for the faculty)

Date	Evaluator	Comments	Score

## CHECKLIST

SNo	Item	Status
1.	Source code documented	Completed
2	Source code uploaded to CCBd server	Completed
3.	Source code in GitLab. Please do not upload your source code to github where it can be seen by everyone.	Completed